

Parallélisme Cuda - TD1

Elana Courtines

2022-09-12

1 Travaux Dirigés de Programmation CUDA

Exercice 1.

Fonction CPU :

```
1 #define BLOCK_SIZE 1024
2
3 void vecADD(float *A, float *B, float *C, int n){
4     // determining the amount of data to declare
5     int bytes = n * sizeof(float);
6     int num_block = (n - 1 + BLOCK_SIZE) / BLOCK_SIZE;
7
8     // create blocks and grid dimension
9     dim3 grid_size = (num_block, 1, 1);
10    dim3 bsize = (BLOCK_SIZE, 1, 1);
11
12    // create the new variables
13    int *dA; int *dB; int *dC;
14
15    // allocate the required memory space
16    cudaMalloc((void **)&dA, bytes);
17    cudaMalloc((void **)&dB, bytes);
18    cudaMalloc((void **)&dC, bytes);
19
20    // copy the given data to the variables used by the GPU
21    cudaMemcpy(dA, A, bytes, cudaMemcpyHostToDevice);
22    cudaMemcpy(dB, B, bytes, cudaMemcpyHostToDevice);
23
24    // launch the GPU function
25    vecAddKernel<<<grid_size,bsize>>>(dA,dB,dC,n);
26
27    // copy the result to the given pointer
28    cudaMemcpy(C, dC, bytes, cudaMemcpyDeviceToHost);
29    cudaFree(dA); cudaFree(dB); cudaFree(dC);
30 }
```

Fonction GPU :

```
1 __global__
2 void vecAddKernel(float *dA, float *dB, float *dC, int n){
3     int indice = blockIdx.x * blockDim.x + threadIdx.x;
4     if (indice < n)
5         dC[indice] = dA[indice] + dB[indice];
6 }
```

Exercice 2 :

Fonction CPU (donnée) :

```
1 #define BLUR_SIZE 3
2 #define BLOCK_SIZE 32
3
4 void blur(unsigned char *in, unsigned char *out, int width, int height){
5     int numbw = (width + BLOCK_SIZE - 1) / BLOCK_SIZE;
6     int numbh = (height + BLOCK_SIZE - 1) / BLOCK_SIZE;
7
8     int bytes = width * height * sizeof(unsigned char);
9     dim3 grid_size = (numbw, numbh, 1);
10    dim3 bsize = (BLOCK_SIZE, BLOCK_SIZE, 1);
11
12    unsigned char *din;
13    unsigned char *dout;
14
15    cudaMalloc((void **)&din, bytes);
16    cudaMalloc((void **)&dout, bytes);
17
18    cudaMemcpy(din, in, bytes, cudaMemcpyHostToDevice);
19
20    blurKernel<<<gdim,bdim>>>(din, dout, width, height);
21
22    cudaMemcpy(out, dout, bytes, cudaMemcpyDeviceToHost);
23    cudaFree(din); cudaFree(dout);
24 }
```

Fonction GPU :

```
1 __global__
2 void blurKernel(unsigned char *din, unsigned char *dout, int width, int height){
3     int lig = blockIdx.y * blockDim.y + threadIdx.y;
4     int col = blockIdx.x * blockDim.x + threadIdx.x;
5
6     if ((lig<height) && (col<width)){ // the pixel must be wihtin the frame
7         int res = 0; // temporary res for the sum
8         int nb = 0; // temporary variable to count the amount
9         of pixels summed
10        // iterate through the width and the height of the Blur Frame
11        for (int currLig = lig-BLUR_SIZE; currLig < lig+BLUR_SIZE+1; currLig++){
12            for (int currCol = col-BLUR_SIZE; currCol < col+BLUR_SIZE+1; currCol++){
13                if ((currLig >= 0) && (currLig < height) && (currCol >= 0) &&
14                (currCol < width)){
15                    res += din[currLig*width+currCol];
16                    nb++;
17                }
18            }
19        }
20        dout[lig*width + col] = (unsigned char) (res / nb);
21    }
22 }
```

Parallélisme Cuda - TD2

Elana Courtines

2022-09-19

Exercice 3 :

Fonction CPU (donnée) :

```
1 void reduce(float *vec, float *sum, int size){
2     float *d_vec;
3     int bytes = size * sizeof(float);
4
5     cudaMalloc((void **)&d_vec, bytes);
6
7     cudaMemcpy(d_vec, vec, bytes, cudaMemcpyHostToDevice);
8
9     kreduce<<<1,size>>>(d_vec, size);
10
11     cudaMemcpy(sum, d_vec, bytes, cudaMemcpyDeviceToHost);
12     cudaFree(d_vec);
13 }
```

Fonction GPU v1:

```
1 __global__
2 void kreduce(float *d_vec, int size){
3     unsigned int tid = threadIdx.x;
4     if (tid < size) {
5         for(int offset = 1; offset<size; offset=offset*2) {
6             if ( tid % 2*offset == 0 ) {
7                 d_vec[tid]+= d_vec[tid+offset];
8             }
9             __syncThreads();
10         }
11     }
12 }
```

Analyse de divergence pour la fonction GPU v1:

itération	#threads	#warps
1	512	32
2	256	32
3	128	32
4	64	32
5	32	32
6	16	16
7	8	8
8	4	4
9	2	2
10	1	1

Fonction GPU v2:

```

1  __global__
2  void kreducev2(float *d_vec, int size){
3      unsigned int tid = threadIdx.x;
4      if (tid < size) {
5          for(int offset = size/2; offset>0; offset=offset/2) {
6              if ( tid < offset ) {
7                  d_vec[tid]+= d_vec[tid+offset];
8              }
9              __syncThreads();
10         }
11     }
12 }

```

Analyse de divergence pour la fonction GPU v2 :

itération (offset)	#threads	#warps
1 (512)	512	16
2 (256)	256	8
3 (128)	128	4
4 (64)	64	2
5 (32)	32	1
6 (16)	16	1
7 (8)	8	1
8 (4)	4	1
9 (2)	2	1
10 (1)	1	1

Exercice 4 :

Fonction CPU :

```
1  #define BLOCK_SIZE 1024
2  #define RADIUS 3
3
4  void convolution(float *in, float *out, float *weight, int size){
5      int num_block = (n - 1 + BLOCK_SIZE) / BLOCK_SIZE;
6      int bytes = size * sizeof(float);
7      float *din;
8      float *dout;
9      cudaMalloc((void **)&din, bytes);
10     cudaMalloc((void **)&dout, bytes);
11     cudaMemcpy(din, in, bytes, cudaMemcpyHostToDevice);
12
13     __constant__ float dweight[2*RADIUS+1];
14     cudaMemcpyToSymbol(dweight, weight, (2*RADIUS+1)*sizeof(float));
15
16     convKernel<<<num_block,BLOCK_SIZE>>>(din, dout, size);
17
18     cudaMemcpy(out, dout, bytes, cudaMemcpyDeviceToHost);
19     cudaFree(din); cudaFree(dout);
20 }
```

Fonction GPU :

```
1 __global__
2 void convKernel(float *in, float *out, int size){
3     int gid = blockIdx.x * blockDim.x + threadIdx.x
4     int tid = threadIdx.x;
5     __shared__ float *sh_in[BLOCK_SIZE + 2*RADIUS];
6     if ( gid < size ){ // start by copying the aligned ones
7         // copy in the cell shifted by 1 #Radius
8         sh_in[tid+RADIUS] = in[gid];
9         // one of the #Radius first threads of the block
10        if (tid < RADIUS) {
11            if ( gid >= RADIUS ){
12                // copy in the cell shifted by 0 #Radius
13                sh_in[tid] = in[gid - RADIUS]
14            } else {
15                sh_in[tid] = 0;
16            }
17        }
18        // one of the #Radius last threads of the block
19        if (tid > BLOCK_SIZE - RADIUS) {
20            if ( gid + BLOCK_SIZE < size ){
21                // copy in the cell shifted by 2 #Radius
22                sh_in[tid + RADIUS*2] = in[gid+RADIUS];
23            } else {
24                sh_in[tid + RADIUS*2] = 0;
25            }
26        }
27    }
28    // The section below has not been corrected in TD
29    __syncthreads();
30    int res=0;
31    for (int i=0; i<2*RADIUS+1; i++){
32        res = res + sh_in[tid-RADIUS+i] * dweight[i];
33    }
34    out[tid]=res;
35 }
```

Parallélisme Cuda - TD3

Elana Courtines

2022-09-26

Exercice 5 :

Fonction CPU :

```
1 void fusion(int *vect1, int *vect2, int size){
2     int bytes = size * sizeof(int);
3     int *d_vin1;
4     int *d_vin2;
5     int *d_vout;
6
7     cudaMalloc((void **)&d_vin1, bytes);
8     cudaMalloc((void **)&d_vin2, bytes);
9     cudaMalloc((void **)&d_vout, 2*bytes);
10
11     cudaMemcpy(d_vin1, in, bytes, cudaMemcpyHostToDevice);
12     cudaMemcpy(d_vin2, in, bytes, cudaMemcpyHostToDevice);
13     cudaMemcpy(d_vout, in, bytes, cudaMemcpyHostToDevice);
14
15     fusionKernel<<<1,size>>>(d_vin1, d_vin2, d_vout, size);
16
17     cudaMemcpy(out, d_vout, bytes, cudaMemcpyDeviceToHost);
18     cudaFree(d_vin1); cudaFree(d_vin2); cudaFree(d_vout);
19 }
```


Fonction GPU :

```
1 __global__
2 void fusionKernel(int *vect1, int *vect2, int *res, int size){
3     int tid = threadIdx.x;
4     __shared__ int *sh_vin1[BLOCK_SIZE]; // it is forbidden to use variables to
5     __shared__ int *sh_vin2[BLOCK_SIZE]; // declare in shared memory
6     int val1;
7     int val2;
8     int pos1;
9     int pos2;
10    if ( tid < size ){
11        sh_vin1[tid] = vect1[tid];
12        sh_vin2[tid] = vect2[tid];
13    }
14    __syncthreads();
15    if ( tid < size ){
16        val1 = sh_vin1[tid];
17        val2 = sh_vin2[tid];
18        pos1 = tid + search(val1, sh_vin2, size)
19        pos2 = tid + search(val2, sh_vin1, size)
20        res[pos1] = val1;
21        res[pos2] = val2;
22    }
23 }
```

Exercice 6 :

Fonction GPU :

```
1 __global__
2 void computeKernel(void *d_bodies, void *d_accel){
3     float4 *bodies = (float4*) d_bodies;
4     float3 *accel = (float3*) d_aceel;
5
6     int tid = threadIdx.x;
7     int gidx = blockIdx.x * blockDim.x + tid;
8
9     __shared__ float4 *sh_bodies[TILE_SIZE];
10
11    float4 mybody = bodies[gidx];
12    float3 acc = {0f, 0f, 0f};
13
14
15    for (int tile=0; tile<NB_BODIES/TILE_SIZE; tile++){
16        int idx = tile * blockDim.x + tid;
17        sh_bodies[tid] = bodies[idx];
18        __syncthreads();
19        for (int k=0; k<TILE_SIZE; k++){
20            interaction(&mybody, &(sh_bodies[k]), &acc); //does the sum for us
21        }
22        __syncthreads();
23    }
24    accel[gidx] = acc;
25 }
```